# RONAK AGARWAL

Senior Data Engineer | 8+ Years of Experience

Pune, Maharashtra | (+91) 8290119822 | agarwalronak60@gmail.com | linkedin.com/in/ronak-ag

## PROFESSIONAL SUMMARY

Senior Data Engineer with 8+ years of experience designing and delivering enterprise-scale data platforms on Azure. Specialized in building real-time and batch ETL pipelines, cloud migrations, and scalable data warehouse architectures. Proven track record of reducing pipeline latency by up to 70%, cutting runtimes by 80%, and driving measurable cost savings across financial services, insurance, and manufacturing clients. Databricks Certified — including the rare GenAI Engineer Associate credential — with deep expertise in PySpark, Kafka, ADF, and Synapse.

## TECHNICAL SKILLS

**Cloud & Big Data:** Azure Synapse, Azure Data Factory (ADF), Azure Databricks, Azure Logic Apps, Azure Function Apps, Blob Storage, Azure SQL DB, Cosmos DB

**Processing & Streaming:** Apache Kafka, PySpark, Apache Spark, Delta Lake, Unity Catalog

**Data Warehousing:** Snowflake, SAP BW, SAP HANA, Matillion, Oracle

**Languages:** Python, SQL , Java , Shell Scripting, ABAP

**DevOps & Governance:** Azure DevOps CI/CD, Git, Databricks Unity Catalog, Jira

**Concepts:** Data Modeling, SCD Implementation, Orchestration Frameworks, System Design, Data Structures & Algorithms

## CERTIFICATIONS

**Databricks Certified Generative AI Engineer Associate** — Databricks *Sept 2025 – Sept 2027*

**Databricks Certified Data Engineer Professional** — Databricks *May 2025 – May 2027*

**Microsoft Azure Fundamentals (AZ-900)** — Microsoft *Sept 2020*

**Big Data 101 & Hadoop 101** — IBM *Aug 2023*

## PROFESSIONAL EXPERIENCE

**Senior Software Engineer** | EPAM Systems (Client: Bosch) Mar 2025 – Present

- Architecting IoT and RDBMS data flow pipelines ingesting real-time sensor data from 50+ device endpoints into a unified Azure data platform.
- Designed and deployed a Kafka-based streaming system that reduced end-to-end data latency by 70% while cutting overall infrastructure costs by ~30%.
- Optimizing PySpark transformation jobs, improving processing efficiency and reducing cloud compute costs through partition tuning and caching strategies.
- Collaborating with cross-functional teams to refine data quality standards and align pipeline architecture with Bosch's analytics roadmap.

**Tools:** *PySpark, Kafka, Azure SQL DB, Azure Data Factory, Azure Logic Apps, Azure Event Hub, Jira*

**Senior Associate** | Synechron (Client: EverestRe) Jan 2023 – Mar 2025

- Led offshore data engineering delivery for EverestRe's ERP team — sole engineer responsible for end-to-end pipeline development, infrastructure scalability, and on-time data delivery across 5+ source systems.
- Designed and built a reusable generic data ingestion framework supporting API, DBMS, and flat-file sources, reducing onboarding time for new source systems by ~60%.
- Implemented Databricks Unity Catalog for workspace-level data governance, enabling fine-grained access control and full data lineage tracking across 10+ datasets.
- Engineered PySpark notebooks and stored procedures for complex transformations; built Azure Logic Apps workflows for automated pipeline success/failure alerting.

**Tools:** *Azure Synapse, PySpark, Azure SQL DB, Azure Data Factory, Azure Logic Apps, Azure Databricks, Jira*

**Recognition:** Applause Award — recognized for outstanding client delivery and initiative

**Consultant** | Deloitte (Client: GM Financial) Sept 2022 – Jan 2023

- Built Synapse pipelines to ingest data from 8+ source systems including APIs and RDBMS into a centralized Azure data warehouse.
- Designed and implemented an orchestration framework for end-to-end job sequencing, reducing pipeline latency by 40% and improving overall throughput.
- Developed SCD (Slowly Changing Dimension) Type 1 & 2 implementations using PySpark notebooks for dimension table management.
- Built CI/CD DevOps pipelines for automated deployment of ADF pipelines and Databricks notebooks across dev, UAT, and production environments.

**Tools:** *Azure Synapse, PySpark, Azure SQL DB, Azure DevOps*

**Recognition:** RR Award — outstanding performance in client delivery

**Consultant** | Deloitte (Client: Solidigm) Apr 2022 – Aug 2022

- Designed and built a multi-source data ingestion pipeline supporting API, DBMS, SAP HANA, and SFTP — covering 12+ distinct data feeds.
- Developed a Python-based Azure Function App to flatten nested JSON/XML files, reducing transformation runtime by 80% versus prior SQL-based approach.
- Designed an end-to-end Audit Batch Log (ABL) framework providing full pipeline observability — tracking row counts, statuses, and error messages for every job run.
- Engaged directly with client architects to align technical design with enterprise data governance and security standards.

**Tools:** *Azure Data Factory, Azure Function App, Snowflake, PySpark*

**Recognition:** RR Award — outstanding performance

**Consultant** | Deloitte (Client: Olympus – PLM) Oct 2021 – Apr 2022

- Designed the end-to-end data flow model covering resource allocation, data retention policies, and network security controls for a new PLM data platform.
- Built ADF orchestration pipelines for ingestion and transformation of PLM data; developed transformation models in Matillion.

**Tools:** *Azure Data Factory, Matillion, Azure DevOps CI/CD, SQL*

**Consultant** | Deloitte (Client: Olympus – Omni) May 2021 – Sept 2022

- Built PySpark transformation notebooks in Azure Databricks for multi-source data processing supporting the Omni channel analytics platform.
- Developed ADF pipelines to ingest data from SharePoint/Microsoft Teams paths with automated email notifications via Logic Apps.
- Created Azure DevOps CI/CD pipelines for ADF and Databricks notebook deployments, eliminating manual release processes and reducing deployment time by ~70%.

**Tools:** *Azure Data Factory, Azure Databricks, PySpark, Python, SQL, Azure DevOps CI/CD, Azure Logic Apps*

**Recognition:** RR Award — outstanding performance

**Systems Engineer → Senior Systems Engineer** | Infosys (Client: ExxonMobil)Dec 2017 – Apr 2021

- Built and maintained ADF pipelines for SAP-to-Azure Blob Storage data migration, supporting ExxonMobil's enterprise data lake ingestion layer.
- Developed SAP BW data flow pipelines and ABAP programs to automate daily SLT monitoring reports, reducing manual effort for ops team.
- Built stored procedures for pipeline log maintenance, enabling end-to-end audit trail across 20+ daily pipeline runs.

**Tools:** *Azure Data Factory, SAP BW, SAP ABAP, SQL DB, Azure Blob Storage*

**Recognition:** Insta Award — recognized for performance and reliability

## EDUCATION

**B. Tech in Computer Science** | Poornima College of Engineering, Jaipur *Apr 2013 – Apr 2017*